
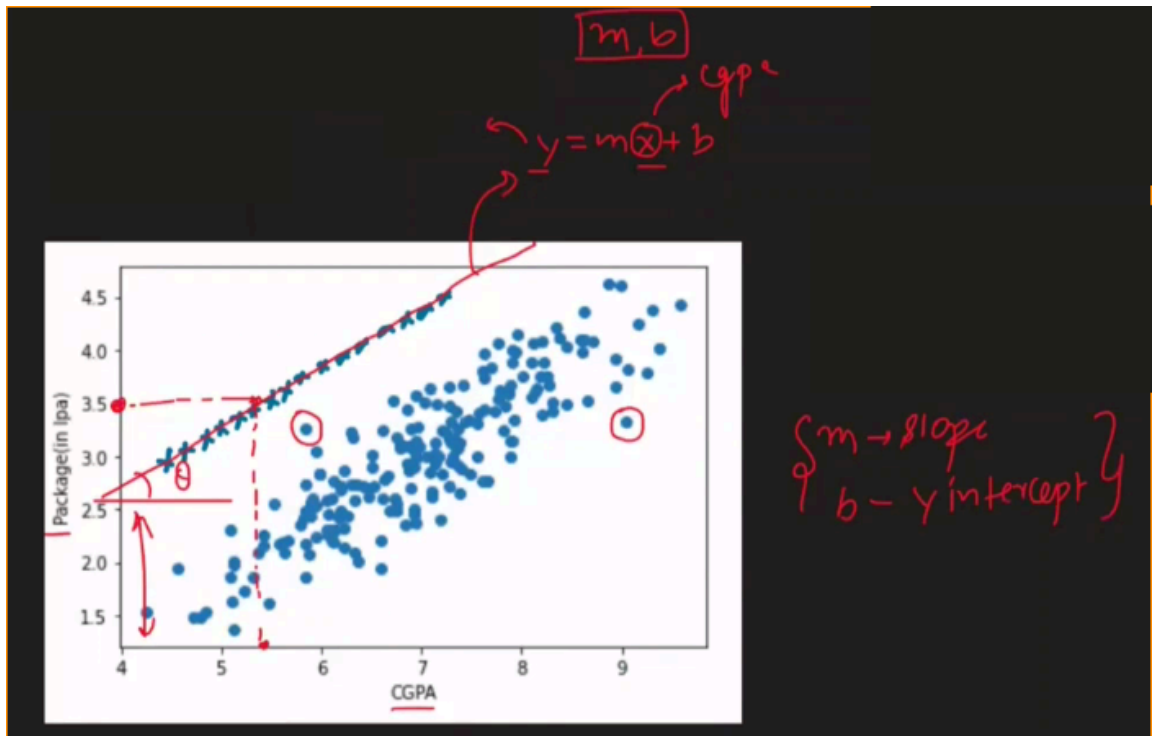



## What is Linear Regression - Geometric Intuition


- SLR is a supervised machine learning algorithm
- Linear Regression generally works when data is linear means it has linear relationship
- There are 3 types:
  1. SLR - When there are only 1 input & 1 output column
  2. MLR - When there are multiple input columns
  3. Polynomial LR - When data is not linear
- suppose we were asked a question: What is the avg package in your college?
- We can simply take the average of packages & quote that amount but it will not represent the true average package bcoz student from different branch & grade will have difference in thier package.
- lets say we have a data contains CGPA & Package so we can plot that data:



- The data shows a general trend, but it's not perfectly straight due to real-world factors we can't precisely measure, like how well someone performed in an interview or a company's urgent need to hire. These uncertainties are called stochastic errors which is why data is sort of linear.
- If our data would have completely linear then we can simply plot a line using equation:  $y = mx + b$
- where  $m$  is the  $\theta$  (angle) &  $b$  is the  $y$  intercept
- 1 unit of change in  $x$  leads to how many units change in  $y$  is " $m$ " or theta
- " $c$ " is the point where best fit line intercepting  $y$  when  $x = 0$



- when new CGPA data comes, we can simply join y to x axis & predict.. but we don't have completely linear data, we have sort of linear data
- we can still draw a line which is called "Best Fit Line"



- **Best fit** line bcoz it has minimum error or trying to pass very closely possible from every point, **Perfect line** is the one which passes through all pont.
- and Thats what **Linear Regression** does, It draws a best fit line on linear data means it calculates the value of those m & b for which line will pass very closely from all points.

## Linear Regression Python Application

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv("placement.csv")
```

```
In [3]: df.sample(7)
```

```
Out[3]:
```


|     | cgpa | package |
|-----|------|---------|
| 66  | 5.11 | 1.63    |
| 111 | 5.42 | 2.25    |
| 167 | 8.13 | 3.60    |
| 27  | 5.42 | 2.16    |
| 78  | 6.59 | 2.21    |
| 75  | 6.97 | 3.28    |
| 105 | 6.66 | 2.82    |

```
In [4]: df.shape
```

```
Out[4]: (200, 2)
```

```
In [5]: #we have 200 student data of cgpa & package..
```

```
In [6]: plt.scatter(df['cgpa'], df['package'])  
plt.xlabel('cgpa')  
plt.ylabel('package')  
plt.show()
```



```
In [7]: #we can see that data is linear so we can apply linear regression model trained on the data  
#which will draw a best fit line and when given new cgpa value, it will predict package
```

```
In [8]: X = df.iloc[:,0:1]  
y = df.iloc[:,1]
```

```
In [9]: #we divide data 2 parts 1 will be used for training & other we can use to test  
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [10]: #now we can train the model  
from sklearn.linear_model import LinearRegression
```

```
In [11]: lr = LinearRegression()
```

```
In [12]: lr.fit(X_train,y_train)
```


```
Out[12]: LinearRegression
LinearRegression()
```

```
In [13]: lr.predict(X_test.iloc[0].values.reshape(1,1))
```

```
C:\Users\iampr\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\base.py:493: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
```

```
Out[13]: array([3.89111601])
```

```
In [14]: plt.scatter(df['cgpa'], df['package'])
plt.plot(X_train, lr.predict(X_train), color='red')
plt.xlabel('cgpa')
plt.ylabel('package')
plt.show()
```



```
In [15]: #above is the best fit line which our lr model finds which also means it has find the
#best fit line is making least mistakes or passing very close to each point
```

```
In [16]: m = lr.coef_
```

```
In [17]: b = lr.intercept_
```

```
In [18]: x = 8.58
y = (m * x) + b
```


```
print(y)
```

```
[3.89111601]
```


```
In [19]: #even if we dont have that value of cgpa still our model will be able to predict  
x = 100 #imaginary number  
y = (m * x) + b  
print(y)
```

```
[54.89908542]
```


## Further Understanding



- $m$  is the weightage, how much cgpa depends on package, if value of  $m$  will be very less then package will be very less dependant on cgpa & vice versa
- means if value of  $m$  changes, change in package will be less since the theta value is very less



- but if slope is very high then less change in cgpa will have more change in package



- if we initialize the value of  $b$  as 0.. and suppose we get a student of experience = 0 (instead of cgpa for understanding) then package should be 0 however freshers also get some salary which is  $b$ . Its called Offset.
- There are 2 ways to find the value of  $m$  &  $b$ 
  1. Closed form solution: It refers to an explicit and direct formula that provides the exact solution to a problem.
  2. Non closed form solution: In non closed form solution we use approximation techniques (calculus, differentiation) to reach the solution.
- **OLS method** is used to find the value of  $m$  &  $b$  using Closed form techniques.
- **Gradient Descent** is used to find the value of  $m$  &  $b$  using non closed form techniques like differentiation.
- Reason why we have 2 techniques is because of complexity.. we use OLS when we have less data & gradient descent when our data is huge.
- Linear Regression class of Scikit learn has OLS technique implemented & SGD regressor class has gradient descent implemented.

## Regression Metrics

- We have multiple metrics bcoz each one have some advantages & disadvantages. Eac will be good on certain type of data
  1. MAE
  2. MSE
  3. RMSE
  4.  $r^2$  score
  5. Adjusted  $r^2$  score

## MAE

- MAE : The average absolute difference between predicted and actual values.
- It is a non-negative value, where lower MAE indicates better accuracy of predictions.

$$\frac{|y_1 - \hat{y}_1| + |y_2 - \hat{y}_2| + \dots + |y_n - \hat{y}_n|}{n}$$

$$mae = \frac{\sum_{i=1}^n |y_i - \hat{y}_i|}{n}$$

Suppose we have actual values  $y = [10, 20, 30, 40]$  and predicted values  $\hat{y} = [12, 18, 28, 38]$ .

1. Calculate the absolute differences:

- $|10 - 12| = 2$
- $|20 - 18| = 2$
- $|30 - 28| = 2$
- $|40 - 38| = 2$

2. Sum these absolute differences:

$$2 + 2 + 2 + 2 = 8$$

3. Calculate the MAE:

$$MAE = \frac{1}{4} \times 8 = 2$$

Therefore, the MAE for this example is 2.

In [20]: `#advantage`



```
# MAE is the loss.. we'll get a number & our goal is to minimize the number  
# MAE has the same unit as output column  
# Robust to outliers
```


In [21]: `#dis-advantage`

```
# MOD is not differentiable at origin & we cant apply optimization techniques like grad  
# MSE(mean squared error ) solves this problem
```

## MSE

- We remove mod & use square instead..

$$mse = \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{n}$$



- **Disadvantage**
  - we'll get a number from MSE but its unit is not same as output, instead it will be square of output.
  - prone to outlier since its taking a square

- **Advantage:** it can be used as a Loss Function since its differentiable

## RMSE

- Its nothing but root of MSE

$$R_{mse} = \sqrt{mse}$$


$$= \sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{n}}$$

**Advantage:** It can be used as a Loss Function since its differentiable & its output is in same unit as output

**disadvantage:** Not robust to outliers

## r2 score

- In r2 score we find how better is our prediction than mean
- Its also coefficient of determination/ goodness of fit



- If r2 score is 0 means SSR/SSM is 1 means SSR & SSM are making same mistakes.
- If r2 score is 1 means SSR/SSM is 0, it can only be 0 when numerator(regression line) is not making any error means its passing through all data points, its a perfect line.
- So after we calculate r2 score, we have to move towards 1 rather than 0

- There are possibilities when  $r^2$  score is -ve, means  $SSR/SSM > 1$  that can only happen when  $SSR > SSM$  means regression is making even more mistakes than mean... that's a worst model.
- If  $r^2$  score is 0.8 --> It means cgpa is able to explain 80% variation in package
- **Problem**
- $r^2$  score increases when we add more input columns even if it's irrelevant like CGPA & Temperature where temperature is completely out of context but  $r^2$  score may increase & to solve this problem we can use adjusted  $r^2$  score

### Adjusted $r^2$ score

Adjusted  $R^2$

$$R^2_{adj} = 1 - \left[ \frac{(1 - R^2)(n - 1)}{(n - 1 - k)} \right]$$

where:

- $r^2$  =  $r^2$  score
- $n$  = no of rows
- $k$  = total no of input cols
- let's assume we add an irrelevant col like temperature
- $k$  (total no of cols) will increase & it will decrease the denominator  $(n - 1 - k)$
- In numerator:  $(n - 1)$  will remain constant since no changes done in no of rows
- $r^2$  score when added irrelevant col:
  - it will remain unchanged: which will make numerator to be remain constant
  - since denominator is decreasing, it will increase whole term inside bracket
  - when we subtract it with 1, then adjusted  $r^2$  score will decrease
- $r^2$  score when added relevant col:
  - $k$  (total no of cols) will increase & it will decrease the denominator  $(n - 1 - k)$
  - $(n - 1)$  is constant
  - $(1 - r^2)$  will decrease faster than denominator since  $r^2$  score will increase faster
  - adjusted  $r^2$  score will increase

- when we are dealing with multiple columns then its better to count on adjusted r2 score.

## Regression metrics Python application

```
In [22]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [23]: y_pred = lr.predict(X_test)
```

```
In [24]: y_test.values
```

```
Out[24]: array([4.1 , 3.49, 2.08, 2.33, 1.94, 1.48, 1.86, 3.09, 4.21, 2.87, 3.65,
                4.   , 2.89, 2.6 , 2.99, 3.25, 1.86, 3.67, 2.37, 3.42, 2.48, 3.65,
                2.6 , 2.83, 4.08, 2.56, 3.58, 3.81, 4.09, 2.01, 3.63, 2.92, 3.51,
                1.94, 2.21, 3.34, 3.34, 3.23, 2.01, 2.61])
```

```
In [25]: print("MAE: ", mean_absolute_error(y_test, y_pred))
          print("MSE: ", mean_squared_error(y_test, y_pred))
          print("RMSE: ", np.sqrt(mean_squared_error(y_test, y_pred)))
          print("R2 Score: ", r2_score(y_test, y_pred))
```

MAE: 0.2884710931878175

MSE: 0.12129235313495527

RMSE: 0.34827051717731616

R2 Score: 0.780730147510384

```
In [26]: r2 = r2_score(y_test, y_pred)
          n = len(y_test)
          k = 1
          r2_adj = 1 - (1 - r2) * (n - 1) / (n - k - 1)
          print("Adjusted R2 Score: ", r2_adj)
```

Adjusted R2 Score: 0.7749598882343415

```
In [ ]:
```